


Application Layer

- Application layer protocols define
 - what messages are exchanged
 - the **syntax** of the messages
 - the **semantics** of the messages
 - **how/when** to exchange messages
 - note the **interaction** is defined, not the application itself!
- We will examine three network applications and their protocols
 - E-mail (SMTP – simple transport protocol) (*Read up in the text book*)
 - the Web (HTTP - HyperText Transfer Protocol)
 - Domain Name Service (DNS)
 - P2P file sharing
- But first we'll look at current application architectures...


Reading:
Kurose & Ross: Ch 2

Computer Networking and Applications



Models of Interaction

- Client - Server
 - central storage of information is always on server
 - distinction between client which receives service and server which provides service
 - note that it is possible for a host to act as both a client and as a server in different interactions.
 - Web, e-mail, FTP
- Peer to Peer
 - distributed storage of information
 - no clear distinction between clients and servers. Hosts share typically equal control of processing and data
 - Peers dynamically join and leave
 - Bit Torrent



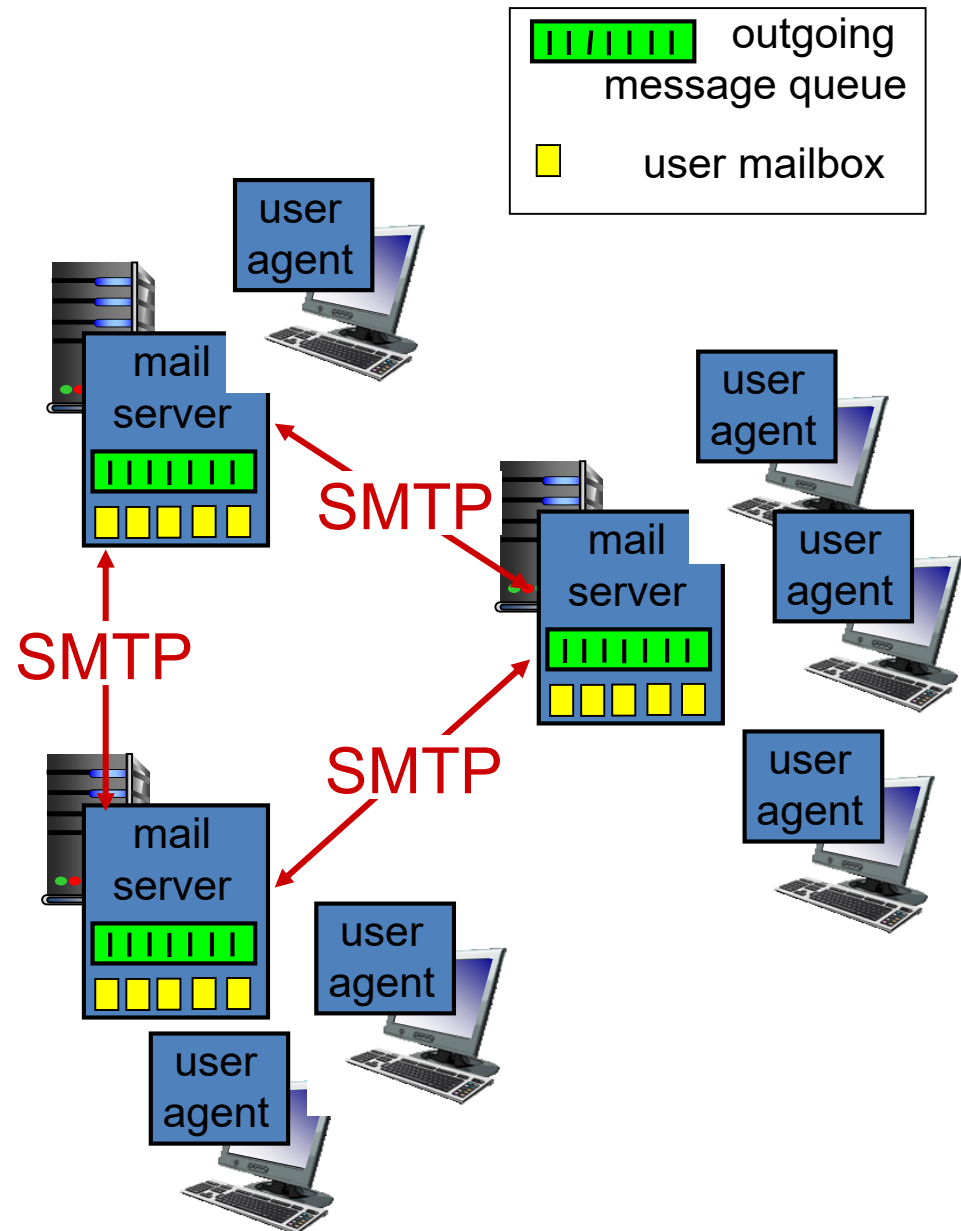
Electronic mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

User Agent


- a.k.a. “mail reader”
- composing, editing, reading mail messages
- e.g., Outlook, Thunderbird, iPhone mail client
- outgoing, incoming messages stored on server



Electronic mail: mail servers

mail servers:

- *mailbox* contains incoming messages for user
- *message queue* of outgoing (to be sent) mail messages
- *SMTP protocol* between mail servers to send email messages
 - client: sending mail server
 - “server”: receiving mail server. TCP port 25



Mail message format


SMTP: protocol for exchanging email msgs

RFC 822: standard for text message format:

- header lines, e.g.,
 - To:
 - From:
 - Subject:

different from SMTP MAIL
FROM, RCPT TO:
commands!

- Body: the “message”
 - ASCII characters only



Sample smtp interaction

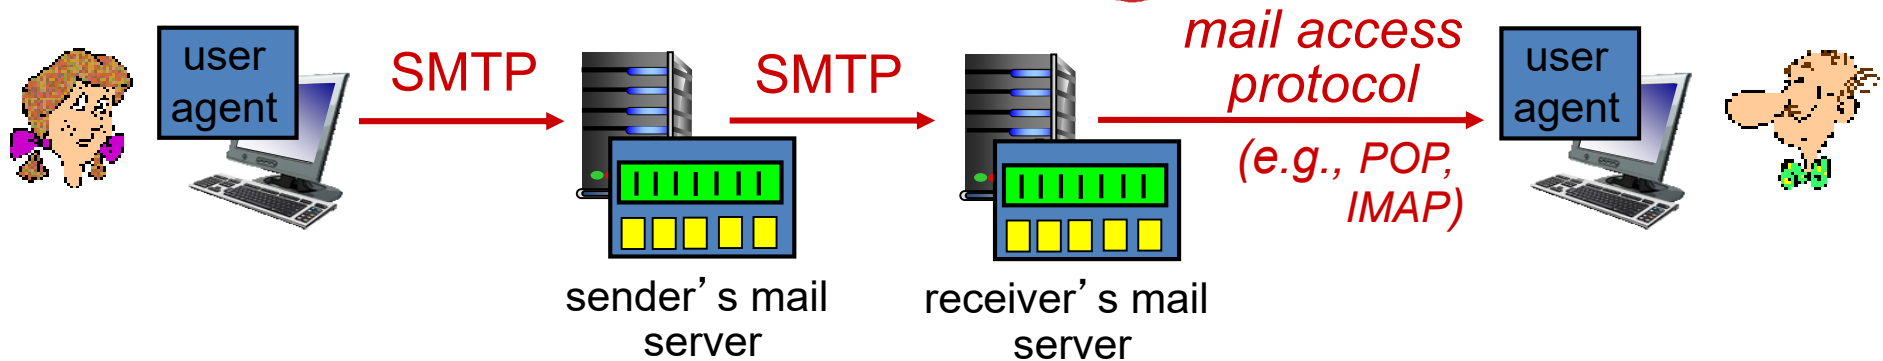
```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C:   How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

h
a
n
d
s
h
a
k
i
n
g

m
e
s
s
a
g
e

c
l
o
s
e

Mail access protocols

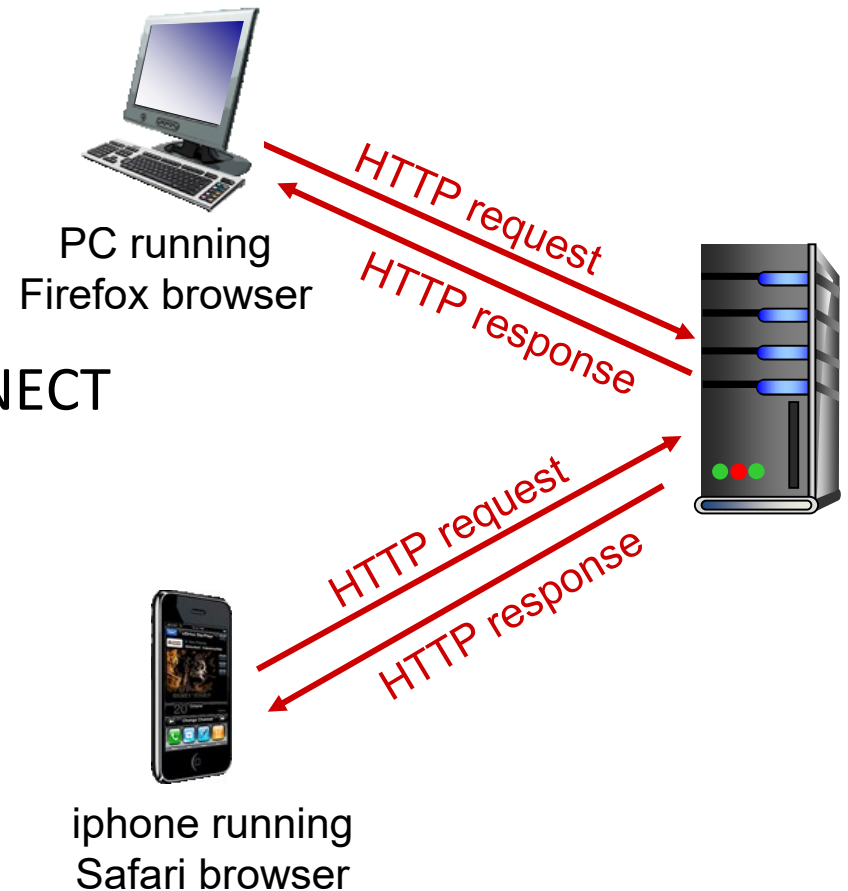


- **SMTP:** delivery/storage to receiver's server
- mail access protocol: retrieval from server
 - **POP:** Post Office Protocol [RFC 1939]: authorization, download
 - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored msgs on server
 - **HTTP:** gmail, Hotmail, Yahoo! Mail, etc.

The World Wide Web

HTTP 1.0 - RFC 1945
HTTP 1.1 - RFC 2616

- Client-Server
- requests
 - GET (get a URL)
 - HEAD (meta data)
 - POST (send data to server)
 - PUT, OPTIONS,DELETE,TRACE,CONNECT
- responses
 - "nnn message <data>"
 - 1xx Informational
 - 2xx Success
 - 3xx Redirection
 - 4xx Client Error
 - 5xx Server Error



HTTP request format


```
GET /somedir/page.html HTTP/1.0
User-agent: Mozilla/4.0
Accept: text/html, image/gif, image/jpeg
Accept-language: fr
```

request line (GET, etc commands)

header lines

(extra carriage return, line feed)

Carriage return, line feed
indicates end of message



HTTP response format

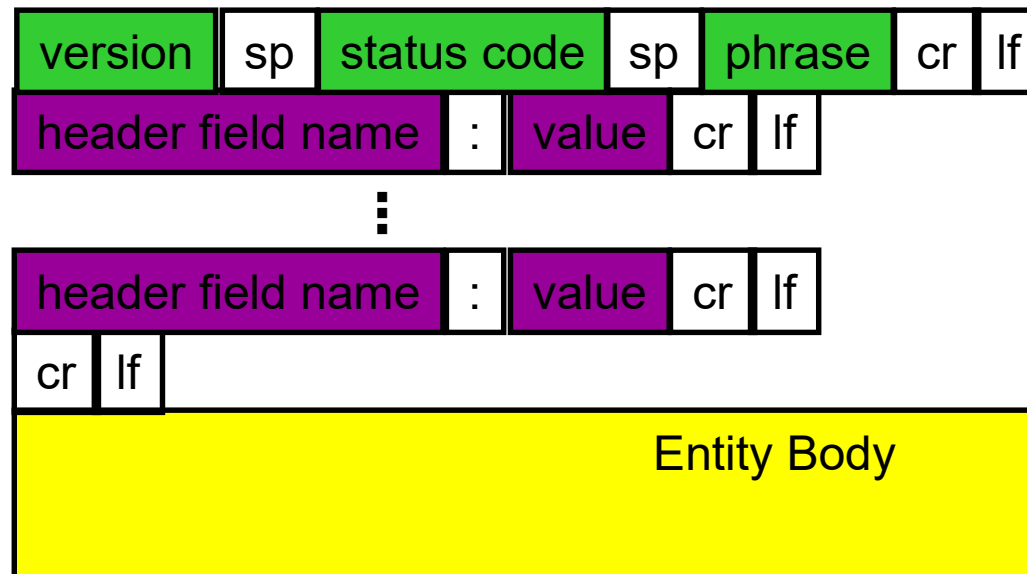
status line
(protocol
status code
status phrase)

header
lines

data, e.g.,
requested
html file

```
HTTP/1.0 200 OK
Date: Thu, 06 Aug 1998 12:00:15 GMT
Server: Apache/1.3.0 (Unix)
Last-Modified: Mon, 22 Jun 1998 .....
Content-Length: 6821
Content-Type: text/html

data data data data data ...
```



Computer Networking and Applications

Trying out network applications (client side) for yourself

1. Telnet to your favorite server (web server in this example):

```
telnet www.cs.adelaide.edu.au 80
```

Opens TCP connection to port 80 (default HTTP server port) at www.cs.adelaide.edu.au
Anything typed in sent to port 80 at www.cs.adelaide.edu.au

2. Type in the protocol request (a GET HTTP request):

```
GET /index.html HTTP/1.1  
Host: www.cs.adelaide.edu.au
```

By typing this in (hit carriage return twice), you send this minimal (but complete) GET request to HTTP server

3. Look at response message sent by server!

You can do the same with a mail server on port 25.

Improving performance - Persistent HTTP

Nonpersistent HTTP issues:

- HTTP/1.0
- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection

Persistent **without** pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

Persistent **with** pipelining:

- default in HTTP/1.1
- Often not turned on in browsers
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

SMTP: final words

- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP server uses CRLF.CRLF to determine end of message


comparison with HTTP:

- HTTP: pull
- SMTP: push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response msg
- SMTP: multiple objects sent in multipart msg

Web caches (proxy server)

Goal: satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client
- Cache acts as both client and server
- Typically cache is installed by ISP (university, company, residential ISP)



Why Web caching?

- Reduce response time for client request.
- Reduce traffic on an institution's access link.
- Internet dense with caches: enables "poor" content providers to effectively deliver content (but so does P2P file sharing)


Caching example:

assumptions:

- ❖ avg object size: 1 Mbit
- ❖ avg request rate from browsers to origin servers: 15/sec
- ❖ **avg data rate to browsers: 15 Mbps**
- ❖ RTT from the *Internet* side of the access link to an origin server: 2 seconds
- ❖ **access link rate: 15 Mbps**

consequences:

- ❖ LAN utilization: 15%
- ❖ access link utilization = **100%** *problem!*
- ❖ total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + usecs




Caching example: fatter access link

assumptions:

- ❖ avg object size: 1 Mbit
- ❖ avg request rate from browsers to origin servers: 15/sec
- ❖ avg data rate to browsers: 15 Mbps
- ❖ RTT from institutional router to any origin server: 2 sec
- ❖ access link rate: ~~15 Mbps~~ → 100 Mbps

consequences:

- ❖ LAN utilization: 15% → 100 Mbps
- ❖ access link utilization = ~~100%~~ → 15%
- ❖ total delay = Internet delay + access delay + LAN delay
= 2 sec + ~~minutes~~ → usecs
msecs




Cost: increased access link speed (not cheap!)

Caching example: install local cache

Calculating access link

utilization, delay with cache:

- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- ❖ access link utilization:
 - 60% of requests use access link
- ❖ data rate to browsers over access link
 - link = $0.6 * 15 \text{ Mbps} = 9 \text{ Mbps}$
 - utilization = $9 \text{ Mbps} / 15 \text{ Mbps} = 60\%$
- ❖ total delay
 - = $0.6 * (\text{delay fetching from origin servers}) + 0.4 * (\text{delay when satisfied at cache or LAN delay})$
 - = $0.6 (2.02) + 0.4 (\sim 10 \text{ ms})$
 - = $\sim 1.2 \text{ secs}$
 - less than with 100 Mbps link (and cheaper too!)



Note: We neglected to LAN delay and access link delay as these are negligible

How do processes (*browser*
and *webserver*) communicate?
(socket programming)

See Sockets.PDF (next week)